

# C++ System Optimizer Project - Interview Q&A;

**Q: Can you explain your C++ System Optimizer project in simple terms?**

A: It's a project that simulates how processes use CPU and memory. I implemented features like dynamic memory allocation, process scheduling, and performance monitoring. The goal was to optimize resource usage and reduce overhead, while also providing profiling to analyze execution efficiency.

**Q: What was the impact of your project?**

A: The optimized memory allocation strategies reduced fragmentation by 30%, and CPU scheduling improved average turnaround time by ~20%. It demonstrated both problem-solving with DSA and system-level optimization skills.

**Q: What C++ features did you use heavily in this project?**

A: I used STL containers (like vector, map, priority\_queue), smart pointers for memory safety, templates for generic code, and multithreading using std::thread. I also applied RAII principles to manage resources.

**Q: How did you ensure memory safety in your project?**

A: By using smart pointers (unique\_ptr and shared\_ptr) instead of raw pointers, RAII patterns, and careful avoidance of memory leaks with profiling tools like Valgrind.

**Q: What is RAII in C++?**

A: RAII stands for Resource Acquisition Is Initialization. It means tying resource management (like memory or file handles) to object lifetime, so resources are released automatically when the object goes out of scope.

**Q: How did you simulate memory management?**

A: I implemented strategies like First-Fit, Best-Fit, and Worst-Fit allocation for dynamic memory blocks, and tracked fragmentation. I compared results across strategies to analyze efficiency.

**Q: Can you explain process scheduling in your system?**

A: Yes, I implemented scheduling algorithms like First-Come-First-Serve (FCFS), Shortest Job First (SJF), and Round Robin. Each had trade-offs in turnaround time, waiting time, and fairness.

**Q: How do you avoid race conditions when using threads in C++?**

A: By using synchronization primitives like mutex, lock\_guard, and condition\_variable. In some cases, I used lock-free structures for performance.

**Q: Can you explain what a lock-free data structure is?**

A: A lock-free structure allows multiple threads to operate on it without traditional locking, usually by leveraging atomic operations (compare-and-swap). This improves performance under high contention.

**Q: What is false sharing in multithreading?**

A: False sharing happens when multiple threads modify variables that reside on the same cache line. It causes unnecessary cache invalidations. It can be avoided by padding or aligning data to cache line boundaries.

**Q: How does cache coherence affect performance?**

A: Cache coherence ensures all cores see a consistent view of memory. Protocols like MESI manage this, but frequent invalidations and cache line bouncing can hurt performance in multithreaded systems.

**Q: What are move semantics in C++ and how did you use them?**

A: Move semantics allow resources to be transferred instead of copied, avoiding expensive deep copies. I used them when passing large containers between functions to reduce overhead.

**Q: Can you explain the difference between compile-time and run-time polymorphism in C++?**

A: Compile-time polymorphism is achieved using templates and function overloading, while run-time polymorphism is achieved using virtual functions. In my project, I used virtual functions for different scheduling algorithms, and templates for generic utilities.

**Q: How do you measure performance in such a project?**

A: By profiling metrics like average turnaround time, waiting time, CPU utilization, and memory fragmentation. Tools like gprof and Valgrind were also helpful.

**Q: What compiler optimizations can you enable in C++ for better performance?**

A: Common flags are -O2 or -O3 for general optimizations, -march=native to use CPU-specific instructions, and -flto for link-time optimizations.

**Q: How does Big-O analysis apply to your project?**

A: Scheduling algorithms had different time complexities (e.g., SJF uses sorting →  $O(n \log n)$ ), and memory allocation strategies varied (First-Fit was  $O(n)$ , Best-Fit could be  $O(n)$  or  $O(n \log n)$  depending on data structure). Optimizing these was part of the project.

**Q: What happens if you delete a pointer twice in C++?**

A: It causes undefined behavior. The second delete might corrupt memory or crash the program. That's why smart pointers are preferred.

**Q: How does virtual memory work?**

A: Virtual memory provides each process with the illusion of having its own large contiguous memory space. The OS maps virtual addresses to physical memory using page tables, and can swap pages to disk if needed.

**Q: Can you explain deadlock and how to prevent it?**

A: Deadlock happens when multiple threads are waiting for each other's resources indefinitely. Prevention strategies include lock ordering, avoiding nested locks, using try\_lock, or employing lock-free techniques.

**Q: How would you scale this project for real-world use?**

A: By integrating more realistic workload models, adding I/O-bound tasks, improving scheduler fairness, and implementing NUMA-aware memory allocation for multi-core systems.